



Manage FlexGroup volumes

ONTAP 9

NetApp
March 29, 2023

Table of Contents

- Manage FlexGroup volumes 1
 - Monitor the space usage of a FlexGroup volume 1
 - Increase the size of a FlexGroup volume..... 3
 - Reduce the size of a FlexGroup volume 4
 - Configure FlexGroup volumes to automatically grow and shrink their size 5
 - Delete directories rapidly on the cluster 6
 - Manage client rights to delete directories rapidly..... 8
 - Create qtrees with FlexGroup volumes 10
 - Use quotas for FlexGroup volumes 10
 - Enable storage efficiency on a FlexGroup volume..... 20
 - Protect FlexGroup volumes using Snapshot copies 21
 - Move the constituents of a FlexGroup volume..... 23
 - Use aggregates in FabricPool for existing FlexGroup volumes 25
 - Rebalance FlexGroup volumes 27

Manage FlexGroup volumes

Monitor the space usage of a FlexGroup volume

You can view a FlexGroup volume and its constituents, and monitor the space used by the FlexGroup volume.

About this task

Beginning with ONTAP 9.6, elastic sizing is supported. ONTAP automatically grows a constituent of a FlexGroup volume if it is running out of space by shrinking any other constituent in the FlexGroup volume that has free space by an equivalent amount. Elastic sizing avoids any out-of-space errors that are generated because of one or more FlexGroup constituent volumes running out of space.



Beginning with ONTAP 9.9.1, logical space reporting and enforcement is also available for FlexGroup volumes. For more information, see [Logical space reporting and enforcement for volumes](#).

Step

1. View the space used by the FlexGroup volume and its constituents: `volume show -vserver vs1 -volume-style-extended flexgroup`

```
cluster-2::> volume show -vserver vs1 -volume-style-extended flexgroup
Vserver   Volume      Aggregate    State      Type      Size
Available Used%
-----
vs1       fg1         -            online     RW        500GB
207.5GB   56%
```

```
ccluster-2::> volume show -vserver vs1 -volume-style-extended flexgroup-
constituent
```

Vserver	Volume	Aggregate	State	Type	Size
Available	Used%				
-----	-----	-----	-----	-----	-----
-----	-----				
vs1	fg1__0001	aggr3	online	RW	31.25GB
12.97GB	56%				
vs1	fg1__0002	aggr1	online	RW	31.25GB
12.98GB	56%				
vs1	fg1__0003	aggr1	online	RW	31.25GB
13.00GB	56%				
vs1	fg1__0004	aggr3	online	RW	31.25GB
12.88GB	56%				
vs1	fg1__0005	aggr1	online	RW	31.25GB
13.00GB	56%				
vs1	fg1__0006	aggr3	online	RW	31.25GB
12.97GB	56%				
vs1	fg1__0007	aggr1	online	RW	31.25GB
13.01GB	56%				
vs1	fg1__0008	aggr1	online	RW	31.25GB
13.01GB	56%				
vs1	fg1__0009	aggr3	online	RW	31.25GB
12.88GB	56%				
vs1	fg1__0010	aggr1	online	RW	31.25GB
13.01GB	56%				
vs1	fg1__0011	aggr3	online	RW	31.25GB
12.97GB	56%				
vs1	fg1__0012	aggr1	online	RW	31.25GB
13.01GB	56%				
vs1	fg1__0013	aggr3	online	RW	31.25GB
12.95GB	56%				
vs1	fg1__0014	aggr3	online	RW	31.25GB
12.97GB	56%				
vs1	fg1__0015	aggr3	online	RW	31.25GB
12.88GB	56%				
vs1	fg1__0016	aggr1	online	RW	31.25GB
13.01GB	56%				

16 entries were displayed.

You can use the available space and percentage space used to monitor the space usage of the FlexGroup volume.

Increase the size of a FlexGroup volume

You can increase the size of a FlexGroup volume either by adding more capacity to the existing constituents of the FlexGroup volume or by expanding the FlexGroup volume with new constituents.

What you'll need

Sufficient space must be available in the aggregates.

About this task

If you want to add more space, you can increase the collective size of the FlexGroup volume. Increasing the size of a FlexGroup volume resizes the existing constituents of the FlexGroup volume.

If you want to improve performance, you can expand the FlexGroup volume. You might want to expand a FlexGroup volume and add new constituents in the following situations:

- New nodes have been added to the cluster.
- New aggregates have been created on the existing nodes.
- The existing constituents of the FlexGroup volume have reached the maximum FlexVol size for the hardware, and therefore the FlexGroup volume cannot be resized.

In releases earlier than ONTAP 9.3, you must not expand FlexGroup volumes after a SnapMirror relationship is established. If you expand the source FlexGroup volume after breaking the SnapMirror relationship in releases earlier than ONTAP 9.3, you must perform a baseline transfer to the destination FlexGroup volume once again. Beginning with ONTAP 9.3, you can expand FlexGroup volumes that are in a SnapMirror relationship.

Step

1. Increase the size of the FlexGroup volume by increasing the capacity or performance of the FlexGroup volume, as required:

If you want to increase the...	Then do this...
Capacity of the FlexGroup volume	Resize the constituents of the FlexGroup volume: <pre>volume modify -vserver vs_server_name -volume fg_name -size new_size</pre>

Performance to the FlexGroup volume	<p>Expand the FlexGroup volume by adding new constituents:</p> <pre>volume expand -vserver vservers_name -volume fg_name -aggr-list aggregate name,... [-aggr-list-multiplier constituents_per_aggr]</pre> <p>The default value of the <code>-aggr-list</code> <code>-multiplier</code> parameter is 1.</p> <p>To expand a FlexGroup volume for FabricPool in ONTAP 9.5, any new aggregates used must be FabricPool.</p>
-------------------------------------	--

Whenever possible, you should increase the capacity of a FlexGroup volume. If you must expand a FlexGroup volume, you should add constituents in the same multiples as the constituents of the existing FlexGroup volume to ensure consistent performance. For example, if the existing FlexGroup volume has 16 constituents with eight constituents per node, you can expand the existing FlexGroup volume by 8 or 16 constituents.

Examples

Example of increasing the capacity of the existing constituents

The following example shows how to add 20 TB space to a FlexGroup volume volX:

```
cluster1::> volume modify -vserver svml -volume volX -size +20TB
```

If the FlexGroup volume has 16 constituents, the space of each constituent is increased by 1.25 TB.

Example of improving performance by adding new constituents

The following example shows how to add two more constituents to the FlexGroup volume volX:

```
cluster1::> volume expand -vserver vs1 -volume volX -aggr-list aggr1,aggr2
```

The size of the new constituents is the same as that of the existing constituents.

Reduce the size of a FlexGroup volume

Beginning with ONTAP 9.6, you can resize a FlexGroup volume to a value lower than its current size to free up the unused space from the volume. When you reduce the size of a FlexGroup volume, ONTAP automatically resizes all of the FlexGroup constituents.

Step

1. Check the current FlexGroup volume size: `'volume size -vserver vservers_name -volume fg_name'`
2. Reduce the size of the FlexGroup volume: `volume size -vserver vservers_name -volume fg_name new_size`

When you specify the new size, you can specify either a lower value than the current size or a negative value using the minus sign (-) by which the current size of the FlexGroup volume is reduced.



If automatic shrinking is enabled for the volume (`volume autosize` command), the minimum autosize is set to the new size of the volume.

The following example displays the current volume size for the FlexGroup volume named `volX` and resizes the volume to 10TB:

```
cluster1::> volume size -vserver svm1 -volume volX
(volume size)
vol size: FlexGroup volume 'svm1:volX' has size 15TB.

cluster1::> volume size -vserver svm1 -volume volX 10TB
(volume size)
vol size: FlexGroup volume 'svm1:volX' size set to 10TB.
```

The following example displays the current volume size for the FlexGroup volume named `volX` and reduces the size of the volume by 5TB:

```
cluster1::> volume size -vserver svm1 -volume volX
(volume size)
vol size: FlexGroup volume 'svm1:volX' has size 15TB.

cluster1::> volume size -vserver svm1 -volume volX -5TB
(volume size)
vol size: FlexGroup volume 'svm1:volX' size set to 10TB.
```

Configure FlexGroup volumes to automatically grow and shrink their size

Beginning with ONTAP 9.3, you can configure FlexGroup volumes to automatically grow and shrink according to how much space they currently require.

What you'll need

The FlexGroup volume must be online.

About this task

You can autosize FlexGroup volumes in two modes:

- Increase the size of the volume automatically (`grow` mode)

Automatic growing helps prevent a FlexGroup volume from running out of space, if the aggregate can supply more space. You can configure the maximum size for the volume. The increase is automatically triggered based on the amount of data being written to the volume in relation to the current amount of used space and any thresholds set.

By default, the maximum size a volume can grow to is 120% of the size at which autogrow is enabled. If you need to ensure that the volume can grow to be larger than that, you must set the maximum size for the volume accordingly.

- Shrink the size of the volume automatically (`grow_shrink` mode)

Automatic shrinking prevents a volume from being larger than needed, freeing space in the aggregate for use by other volumes.

Autoshrink can only be used in combination with autogrow to meet changing space demands and is not available alone. When autoshrink is enabled, ONTAP automatically manages the shrinking behavior of a volume to prevent an endless loop of autogrow and autoshrink actions.

As a volume grows, the maximum number of files it can contain might be automatically increased. When a volume is shrunk, the maximum number of files it can contain is left unchanged, and a volume cannot be automatically shrunk below the size that corresponds to its current maximum number of files. For this reason, it might not be possible to automatically shrink a volume all the way to its original size.

Step

1. Configure the volume to grow and shrink its size automatically: `volume autosize -vserver vserver_name -volume vol_name -mode [grow | grow_shrink]`

You can also specify the maximum size, minimum size, and thresholds for growing or shrinking the volume.

The following command enables automatic size changes for a volume called `fg1`. The volume is configured to grow to a maximum size of 5 TB when it is 70% full.

```
cluster1::> volume autosize -volume fg1 -mode grow -maximum-size 5TB
-grow-threshold-percent 70
vol autosize: volume "vs_src:fg1" autosize settings UPDATED.
```

Delete directories rapidly on the cluster

Beginning with ONTAP 9.8, you can use low-latency *fast-directory delete* functionality to delete directories from Linux and Windows client shares asynchronously (that is, in the background). Cluster and SVM administrators can perform asynchronous delete operations on both FlexVol and FlexGroup volumes.

If you are using a version of ONTAP earlier than ONTAP 9.11.1, you must be a cluster administrator or a SVM administrator using the advanced privilege mode.

Beginning with ONTAP 9.11.1, a storage administrator can grant rights on a volume to allow NFS and SMB clients to perform asynchronous delete operations. For more information, see [Manage client rights to delete directories rapidly](#).

Beginning with ONTAP 9.8, you can use fast directory delete functionality using the ONTAP CLI. Beginning with ONTAP 9.9.1, you can use this functionality with System Manager. For more information about this process, see [Take corrective action based on analytics](#).

System Manager

For more information, see [Take corrective action based on analytics](#).

CLI

Use the CLI to perform a fast directory delete

1. Enter advanced privilege mode:

```
-privilege advance
```

2. Delete files or directories on a FlexVol or FlexGroup volume:

```
volume file async-delete start -vserver vs1 -volume vol1  
-path file_path -throttle throttle
```

The minimum throttle value is 10, the maximum is 100,000, and the default is 5000.

The following example deletes the directory named d2, which is located in the directory named d1.

```
cluster::*>volume file async-delete start -vserver vs1 -volume vol1  
-path d1/d2
```

3. Verify that the directory was deleted:

```
event log show
```

The following example shows output for the event log when the directory is successfully deleted.

```
cluster-cli::*> event log show
```

Time	Node	Severity	Event
MM/DD/YYYY 00:11:11	cluster-vsim	INFORMATIONAL	asyncDelete.message.success: Async delete job on path d1/d2 of volume (MSID: 2162149232) was completed.

Cancel a directory delete job

1. Enter advanced privilege mode:

```
set -privilege advanced
```

2. Verify that the directory delete is in progress:

```
volume file async-delete show
```

If the SVM, volume, JobID, and path of your directory is displayed, you can cancel the job.

3. Cancel the directory delete:

```
volume file async-delete cancel -vserver SVM_name -volume volume_name
-jobid job_id
```

Manage client rights to delete directories rapidly

Beginning with ONTAP 9.11.1, storage administrators can grant rights on a volume to allow NFS and SMB clients to perform low latency *fast-directory delete* operations themselves. When asynchronous delete is enabled on the cluster, Linux client users can use the `mv` command and Windows client users can use the `rename` command to delete a directory rapidly on the specified volume by moving it to a hidden directory that by default is named `.ontaptrashbin`.

Enable client asynchronous directory delete

Steps

1. From the cluster CLI, enter advanced privilege mode: `-privilege advance`
2. Enable client asynchronous delete and, if desired, provide an alternate name for the trashbin directory:

```
volume file async-delete client enable volume volname vserver vserverName
trashbinname name
```

Example using the default trashbin name:

```
cluster1::*> volume file async-delete client enable -volume v1 -vserver
vs0
```

```
Info: Async directory delete from the client has been enabled on volume
"v1" in
      Vserver "vs0".
```

Example specifying an alternate trashbin name:

```
cluster1::*> volume file async-delete client enable -volume test
-trashbin ntaptrash -vserver vs1
```

```
Success: Async directory delete from the client is enabled on volume
"v1" in
      Vserver "vs0".
```

3. Verify client asynchronous delete is enabled:

```
volume file async-delete client show
```

Example:

```
cluster1::*> volume file async-delete client show
```

Vserver	Volume	async-delete client	TrashBinName
vs1	vol1	Enabled	.ntaptrash
vs2	vol2	Disabled	-

2 entries were displayed.

Disable client asynchronous directory delete

Steps

1. From the cluster CLI, disable client asynchronous directory delete:

```
volume file async-delete client disable volume volname vservers vserverName
```

Example:

```
cluster1::*> volume file async-delete client disable -volume vol1  
-vserver vs1
```

Success: Asynchronous directory delete client disabled
successfully on volume.

2. Verify client asynchronous delete is disabled:

```
volume file async-delete client show
```

Example:

```
cluster1::*> volume file async-delete client show
```

Vserver	Volume	async-delete client	TrashBinName
vs1	vol1	Disabled	-
vs2	vol2	Disabled	-

2 entries were displayed.

Create qtrees with FlexGroup volumes

Beginning with ONTAP 9.3, you can create qtrees with FlexGroup volumes. Qtrees enable you to partition your FlexGroup volumes into smaller segments that you can manage individually.

About this task

- If you want to revert to ONTAP 9.2 or earlier and if you have created one or more qtrees in the FlexGroup volume or modified the attributes (security style and SMB oplocks) of the default qtree, you must delete all of the non-default qtrees and then disable the qtree functionality on each FlexGroup volume before reverting to ONTAP 9.2 or earlier.

[Disable qtree functionality in FlexGroup volumes before reverting](#)

- If the source FlexGroup volume has qtrees in a SnapMirror relationship, the destination cluster must be running ONTAP 9.3 or later (a version of ONTAP software that supports qtrees).
- Beginning with ONTAP 9.5, qtree statistics are supported for FlexGroup volumes.

Steps

1. Create a qtree in the FlexGroup volume: `volume qtree create -vserver vs0 -volume fg1 -qtree qtree1 -security-style mixed`

You can optionally specify the security style, SMB oplocks, UNIX permissions, and export policy for the qtree.

```
cluster1::> volume qtree create -vserver vs0 -volume fg1 -qtree qtree1  
-security-style mixed
```

Related information

[Logical storage management](#)

Use quotas for FlexGroup volumes

In ONTAP 9.4 and earlier, you can apply quotas rules to FlexGroup volumes only for reporting purposes, but not for enforcing quota limits. Beginning with ONTAP 9.5, you can enforce limits on quota rules that are applied to FlexGroup volumes.

About this task

- Beginning with ONTAP 9.5, you can specify hard, soft, and threshold limit quotas for FlexGroup volumes.

You can specify these limits to constrain the amount of space, the number of files that a specific user, group, or qtree can create, or both. Quota limits generate warning messages in the following scenarios:

- When usage exceeds a configured soft limit, ONTAP issues a warning message, but further traffic is still allowed.

If usage later drops below the configured soft limit again, an all-clear message is issued.

- When usage exceeds a configured threshold limit, ONTAP issues a second warning message.

No all-clear administrative message is issued when usage later drops below a configured threshold limit.

- If usage reaches a configured hard limit, ONTAP prevents further resource consumption by rejecting traffic.
- In ONTAP 9.5, quota rules cannot be created or activated on the destination FlexGroup volume of a SnapMirror relationship.
- During quota initialization, quotas are not enforced, and there are no notifications of breached quotas following quota initialization.

To check if quotas were breached during quota initialization, you can use the `volume quota report` command.

Quota targets and types

Quotas have a type: they can be either user, group, or tree. Quota targets specify the user, group, or qtree for which the quota limits are applied.

The following table lists the kinds of quota targets, what types of quotas each quota target is associated with, and how each quota target is represented:

Quota target	Quota type	How target is represented	Notes
user	user quota	UNIX user name UNIX UID Windows user name in pre-Windows 2000 format Windows SID	User quotas can be applied for a specific volume or qtree.
group	group quota	UNIX group name UNIX GID	Group quotas can be applied for a specific volume or qtree. <div>  <p>ONTAP does not apply group quotas based on Windows IDs.</p> </div>
qtree	tree quota	qtree name	Tree quotas are applied to a particular volume and do not affect qtrees in other volumes.

""	user quota group quota tree quota	Double quotation marks (<code>""</code>)	A quota target of <code>""</code> denotes a <i>default quota</i> . For default quotas, the quota type is determined by the value of the <code>type</code> field.
----	---	---	--

Behavior of FlexGroup volumes when quota limits are exceeded

Beginning with ONTAP 9.5, quota limits are supported on FlexGroup volumes. There are some differences in the way quota limits are enforced on a FlexGroup volume when compared to a FlexVol volume.

FlexGroup volumes might show the following behaviors when the quota limits are exceeded:

- The space and file usage in a FlexGroup volume might reach up to 5 percent higher than the configured hard limit before the quota limit is enforced by rejecting further traffic.

To provide the best performance, ONTAP might allow the space consumption to exceed the configured hard limit by a small margin before the quota enforcement begins. This additional space consumption does not exceed 5 percent of the configured hard limits, 1 GB, or 65536 files, whichever is lower.

- After the quota limit is reached, if a user or administrator deletes some files or directories such that the quota usage is now below the limit, the subsequent quota-consuming file operation might resume with a delay (might take up to 5 seconds to resume).
- When the total space and file usage of a FlexGroup volume exceed the configured quota limits, there might be a slight delay in logging an event log message.
- You might get “no space” errors if some constituents of the FlexGroup volume get full, but the quota limits are not reached.
- Operations, such as renaming a file or directory or moving files between qtrees, on quota targets, for which quota hard limits are configured, might take longer when compared to similar operations on FlexVol volumes.

Examples of quota enforcement for FlexGroup volumes

You can use the examples to understand how to configure quotas with limits in ONTAP 9.5 and later.

Example 1: Enforcing a quota rule with disk limits

1. You should create a quota policy rule of type `user` with both an achievable soft disk limit and hard disk limit.

```
cluster1::> volume quota policy rule create -vserver vs0 -policy-name
default -volume FG -type user -target "" -qtree "" -disk-limit 1T -soft
-disk-limit 800G
```

2. You can view the quota policy rule:

```
cluster1::> volume quota policy rule show -vserver vs0 -policy-name
default -volume FG
```

```
Vserver: vs0                Policy: default                Volume: FG
```

Type	Target	Qtree	User Mapping	Disk Limit	Soft Disk Limit	Files Limit	Soft Files Limit
Threshold							
-----	-----	-----	-----	-----	-----	-----	-----

user	""	""	off	1TB	800GB	-	-
-							

3. To activate the new quota rule, you initialize quotas on the volume:

```
cluster1::> volume quota on -vserver vs0 -volume FG -foreground true
[Job 49] Job succeeded: Successful
```

4. You can view the disk usage and file usage information of the FlexGroup volume by using the quota report.

```
cluster1::> volume quota report -vserver vs0 -volume FG
Vserver: vs0
```

Volume Specifier	Tree	Type	ID	----Disk----	----Files----	Quota
				Used Limit	Used Limit	
-----	-----	-----	-----	-----	-----	-----

FG		user	root	50GB -	1 -	
FG		user	*	800GB 1TB	0 -	*

2 entries were displayed.

After the hard disk limit is reached, the quota policy rule target (user, in this case) is blocked from writing more data to the files.

Example 2: Enforcing a quota rule for multiple users

1. You should create a quota policy rule of type `user`, where multiple users are specified in the quota target (UNIX users, SMB users, or a combination of both) and where the rule has both an achievable soft disk limit and hard disk limit.

```
cluster1::> quota policy rule create -vserver vs0 -policy-name default
-volume FG -type user -target "rdavis,ABCCORP\RobertDavis" -qtree ""
-disk-limit 1TB -soft-disk-limit 800GB
```

2. You can view the quota policy rule:

```
cluster1::> quota policy rule show -vserver vs0 -policy-name default
-volume FG
```

```
Vserver: vs0                Policy: default                Volume: FG
```

Type	Target	Qtree	User Mapping	Disk Limit	Soft Disk Limit	Files Limit	Soft Files Limit
user	"rdavis,ABCCORP\RobertDavis"	""	off	1TB	800GB	-	-

3. To activate the new quota rule, you initialize quotas on the volume:

```
cluster1::> volume quota on -vserver vs0 -volume FG -foreground true
[Job 49] Job succeeded: Successful
```

4. You can verify that the quota state is active:

```
cluster1::> volume quota show -vserver vs0 -volume FG
Vserver Name: vs0
Volume Name: FG
Quota State: on
Scan Status: -
Logging Messages: on
Logging Interval: 1h
Sub Quota Status: none
Last Quota Error Message: -
Collection of Quota Errors: -
```

5. You can view the disk usage and file usage information of the FlexGroup volume by using the quota report.


```
cluster1::> quota report -vserver vs0 -volume FG
Vserver: vs0
```

Volume Specifier	Tree	Type	ID	----Disk----		----Files-----		Quota
				Used	Limit	Used	Limit	
FG		user	rdavis,ABCCORP\RobertDavis	0B	1TB	0	-	

The quota limit is shared among all users listed in the quota target.

After the hard disk limit is reached, users listed in the quota target are blocked from writing more data to the files.

Example 3: Enforcing quota with user mapping enabled

1. You should create a quota policy rule of type user, specify a UNIX user or a Windows user as the quota target with user-mapping set to on, and create the rule with both an achievable soft disk limit and hard disk limit.

The mapping between UNIX and Windows users must be configured earlier by using the `vserver name-mapping create` command.

```
cluster1::> quota policy rule create -vserver vs0 -policy-name default
-volume FG -type user -target rdavis -qtree "" -disk-limit 1TB -soft
-disk-limit 800GB -user-mapping on
```

2. You can view the quota policy rule:

```
cluster1::> quota policy rule show -vserver vs0 -policy-name default
-volume FG
```

Vserver: vs0			Policy: default			Volume: FG	
Type	Target	Qtree	User Mapping	Disk Limit	Soft Disk Limit	Files Limit	Soft Files Limit
user	rdavis	""	on	1TB	800GB	-	-

3. To activate the new quota rule, you initialize quotas on the volume:

```
cluster1::> volume quota on -vserver vs0 -volume FG -foreground true
[Job 49] Job succeeded: Successful
```

4. You can verify that the quota state is active:

```
cluster1::> volume quota show -vserver vs0 -volume FG
      Vserver Name: vs0
      Volume Name: FG
      Quota State: on
      Scan Status: -
      Logging Messages: on
      Logging Interval: 1h
      Sub Quota Status: none
      Last Quota Error Message: -
      Collection of Quota Errors: -
```

5. You can view the disk usage and file usage information of the FlexGroup volume by using the quota report.

```
cluster1::> quota report -vserver vs0 -volume FG
Vserver: vs0
```

Volume	Tree	Type	ID	-----Disk-----	-----Files-----	Quota	
Specifier				Used	Limit	Used	Limit
-----	-----	-----	-----	-----	-----	-----	-----
FG		user	rdavis,ABCCORP\RobertDavis	0B	1TB	0	-
rdavis							

The quota limit is shared between the user listed in the quota target and its corresponding Windows or UNIX user.

After the hard disk limit is reached, both the user listed in the quota target and its corresponding Windows or UNIX user is blocked from writing more data to the files.

Example 4: Verifying the qtree size when quota is enabled

1. You should create a quota policy rule of type `tree` and where the rule has both an achievable soft disk limit and hard disk limit.

```
cluster1::> quota policy rule create -vserver vs0 -policy-name default
-volume FG -type tree -target tree_4118314302 -qtree "" -disk-limit 48GB
-soft-disk-limit 30GB
```

2. You can view the quota policy rule:

```
cluster1::> quota policy rule show -vserver vs0
```

Vserver: vs0			Policy: default			Volume: FG	
			User	Disk	Soft	Files	Soft
Type	Target	Qtree	Mapping	Limit	Disk	Files	Files
Threshold					Limit	Limit	Limit
-----	-----	-----	-----	-----	-----	-----	-----
tree	tree_4118314302	""	-	48GB	-	20	-

3. To activate the new quota rule, you initialize quotas on the volume:

```
cluster1::> volume quota on -vserver vs0 -volume FG -foreground true
[Job 49] Job succeeded: Successful
```

a. You can view the disk usage and file usage information of the FlexGroup volume by using the quota report.

```
cluster1::> quota report -vserver vs0
```

Vserver: vs0									
----Disk----		----Files-----		Quota					
Volume	Tree	Type	ID	Used	Limit	Used	Limit	Specifier	
FG	tree_4118314302	tree	1	30.35GB	48GB	14	20	tree_4118314302	

The quota limit is shared between the user listed in the quota target and its corresponding Windows or UNIX user.

4. From an NFS client, use the `df` command to view the total space usage, available space, and the used space.

```
scsps0472342001# df -m /t/10.53.2.189/FG-3/tree_4118314302
```

Filesystem	1M-blocks	Used	Available	Use%	Mounted on
10.53.2.189/FG-3	49152	31078	18074	63%	/t/10.53.2.189/FG-3

With hard limit, the space usage is calculated from an NFS client as follows:

- Total space usage = hard limit for tree
- Free space = Hard limit minus qtree space usage Without hard limit, the space usage is calculated from an NFS client as follows:
- Space usage = quota usage
- Total space = Sum of quota usage and physical free space in the volume

5. From the SMB share, use Windows Explorer to view the total space usage, available space, and the used space.

From an SMB share, you should be aware of the following considerations for calculating the space usage:

- The user quota hard limit for the user and group is taken into consideration for calculating the total available space.
- The minimum value among the free space of the tree quota rule, the user quota rule, and the group quota rule is considered as the free space for the SMB share.
- The total space usage is variable for SMB and depends on the hard limit that corresponds to the minimum free space among the tree, user, and group.

Apply rules and limits on the FlexGroups volume

Steps

1. Create quota rules for targets : `volume quota policy rule create -vserver vs0 -policy -name quota_policy_of_the_rule -volume flexgroup_vol -type {tree|user|group} -target target_for_rule -qtree qtree_name [-disk-limit hard_disk_limit_size] [-file-limit hard_limit_number_of_files] [-threshold threshold_disk_limit_size] [-soft-disk-limit soft_disk_limit_size] [-soft-file-limit soft_limit_number_of_files]`

- In ONTAP 9.2 and ONTAP 9.1, the quota target type can be only `user` or `group` for FlexGroup volumes.

Tree quota type is not supported for FlexGroup volumes in ONTAP 9.2 and ONTAP 9.1.

- In ONTAP 9.3 and later, the quota target type can be `user`, `group`, or `tree` for FlexGroup volumes.
- A path is not supported as the target when creating quota rules for FlexGroup volumes.
- Beginning with ONTAP 9.5, you can specify hard disk limit, hard file limit, soft disk limit, soft file limit, and threshold limit quotas for FlexGroup volumes.

In ONTAP 9.4 and earlier, you cannot specify the disk limit, file limit, threshold for disk limit, soft disk limit, or soft file limit when you create quota rules for FlexGroup volumes.

The following example shows a default quota rule being created for the user target type:

```
cluster1::> volume quota policy rule create -vserver vs0 -policy-name
quota_policy_vs0_1 -volume fg1 -type user -target "" -qtree ""
```

The following example shows a tree quota rule being created for the qtree named qtree1:

```
cluster1::> volume quota policy rule create -policy-name default -vserver
vs0 -volume fg1 -type tree -target "qtree1"
```

1. Activate the quotas for the specified FlexGroup volume: `volume quota on -vserver svm_name -volume flexgroup_vol -foreground true`

```
cluster1::> volume quota on -vserver vs0 -volume fg1 -foreground true
```

1. Monitor the state of quota initialization: `volume quota show -vserver svm_name`

FlexGroup volumes might show the mixed state, which indicates that all of the constituent volumes are not in the same state yet.

```
cluster1::> volume quota show -vserver vs0
```

Vserver	Volume	State	Scan Status
vs0	fg1	initializing	95%
vs0	vol1	off	-

2 entries were displayed.

1. View the quota report for the FlexGroup volume with active quotas: `volume quota report -vserver svm_name -volume flexgroup_vol`

You cannot specify a path with the `volume quota report` command for FlexGroup volumes.

The following example shows the user quota for the FlexGroup volume fg1:

```
cluster1::> volume quota report -vserver vs0 -volume fg1
```

Vserver: vs0

Quota				----Disk----		----Files----	
Volume	Tree	Type	ID	Used	Limit	Used	Limit
Specifier							
fg1		user	*	0B	-	0	-
fg1		user	root	1GB	-	1	-

2 entries were displayed.

The following example shows the tree quota for the FlexGroup volume fg1:

```
cluster1::> volume quota report -vserver vs0 -volume fg1
Vserver: vs0
```

Volume Specifier	Tree	Type	ID	----Disk----		----Files-----		Quota
				Used	Limit	Used	Limit	
fg1	qtreet1	tree	1	68KB	-	18	-	
qtreet1		tree	*	0B	-	0	-	*

2 entries were displayed.

Results

The quota rules and limits are applied on the FlexGroups volume.

The usage might reach up to 5 percent higher than a configured hard limit before ONTAP enforces the quota by rejecting further traffic.

Related information

[ONTAP 9 Commands](#)

Enable storage efficiency on a FlexGroup volume

You can run deduplication and data compression together or independently on a FlexGroup volume to achieve optimal space savings.

What you'll need

The FlexGroup volume must be online.

Steps

1. Enable storage efficiency on the FlexGroup volume: `volume efficiency on -vserver svm_name -volume volume_name`

Storage efficiency operations are enabled on all the constituents of the FlexGroup volume.

If a FlexGroup volume is expanded after storage efficiency is enabled on the volume, storage efficiency is automatically enabled on the new constituents.

2. Enable the required storage efficiency operation on the FlexGroup volume by using the `volume efficiency modify` command.

You can enable inline deduplication, postprocess deduplication, inline compression, and postprocess compression on FlexGroup volumes. You can also set the type of compression (secondary or adaptive) and specify a schedule or efficiency policy for the FlexGroup volume.

3. If you are not using schedules or efficiency policies for running the storage efficiency operations, start the efficiency operation: `volume efficiency start -vserver svm_name -volume volume_name`

If deduplication and data compression are enabled on a volume, data compression is run initially followed by deduplication. This command fails if any efficiency operation is already active on the FlexGroup volume.

4. Verify the efficiency operations that are enabled on the FlexGroup volume: `volume efficiency show -vserver svm_name -volume volume_name`

```
cluster1::> volume efficiency show -vserver vs1 -volume fg1
      Vserver Name: vs1
      Volume Name: fg1
      Volume Path: /vol/fg1
      State: Enabled
      Status: Idle
      Progress: Idle for 17:07:25
      Type: Regular
      Schedule: sun-sat@0

...

      Compression: true
      Inline Compression: true
      Incompressible Data Detection: false
      Constituent Volume: false
      Compression Quick Check File Size: 524288000
      Inline Dedupe: true
      Data Compaction: false
```

Protect FlexGroup volumes using Snapshot copies

You can create Snapshot policies that automatically manage the creation of Snapshot copies or you can manually create Snapshot copies for FlexGroup volumes. A valid Snapshot copy is created for a FlexGroup volume only after ONTAP can successfully create a Snapshot copy for each constituent of the FlexGroup volume.

About this task

- If you have multiple FlexGroup volumes associated with a Snapshot policy, you should ensure that the FlexGroup volumes schedules do not overlap.
- Beginning with ONTAP 9.8, the maximum number of Snapshot copies supported on a FlexGroup volume is 1023.



Beginning with ONTAP 9.8, the `volume snapshot show` command for FlexGroup volumes reports Snapshot copy size using logical blocks, rather than calculating the youngest owned blocks. This new size calculation method might make the Snapshot copy size appear larger than calculations in earlier versions of ONTAP.

Steps

1. Create a Snapshot policy or manually create a Snapshot copy:

If you want to create a...	Enter this command...
Snapshot policy	<pre>volume snapshot policy create</pre> <div>  <p>The schedules that are associated with the Snapshot policy of a FlexGroup volume must have an interval greater than 30 minutes.</p> </div> <p>When you create a FlexGroup volume, the default Snapshot policy is applied to the FlexGroup volume.</p>
Snapshot copy manually	<pre>volume snapshot create</pre> <div>  <p>After you create a Snapshot copy for a FlexGroup volume, you cannot modify the attributes of the Snapshot copy. If you want to modify the attributes, you must delete and then re-create the Snapshot copy.</p> </div>

Client access to the FlexGroup volume is briefly quiesced when a Snapshot copy is created.

1. Verify that a valid Snapshot copy is created for the FlexGroup volume: `volume snapshot show -volume volume_name -fields state`

```
cluster1::> volume snapshot show -volume fg -fields state
vserver volume snapshot                state
-----
fg_vs    fg        hourly.2016-08-23_0505 valid
```

2. View the Snapshot copies for the constituents of the FlexGroup volume: `volume snapshot show -is -constituent true`


```
cluster1::> volume snapshot show -is-constituent true
```

---Blocks---				
Vserver	Volume	Snapshot	Size	Total%
Used%				
-----	-----	-----	-----	-----
fg_vs	fg__0001	hourly.2016-08-23_0505	72MB	0%
27%				
	fg__0002	hourly.2016-08-23_0505	72MB	0%
27%				
	fg__0003	hourly.2016-08-23_0505	72MB	0%
27%				
...				
	fg__0016	hourly.2016-08-23_0505	72MB	0%
27%				

Move the constituents of a FlexGroup volume

You can move the constituents of a FlexGroup volume from one aggregate to another for balancing the load when certain constituents experience more traffic. Moving constituents also helps in freeing up space on an aggregate for resizing the existing constituents.

What you'll need

To move a FlexGroup volume constituent that is in a SnapMirror relationship, you must have initialized the SnapMirror relationship.

About this task

You cannot perform a volume move operation while the constituents of the FlexGroup volume are being expanded.

Steps

1. Identify the FlexGroup volume constituent that you want to move: `volume show -vserver svm_name -is-constituent *`

```
cluster1::> volume show -vserver vs2 -is-constituent *
```

Vserver	Volume	Aggregate	State	Type	Size
Available	Used%				
vs2	fg1	-	online	RW	400TB
15.12TB	62%				
vs2	fg1__0001	aggr1	online	RW	25TB
8.12MB	59%				
vs2	fg1__0002	aggr2	online	RW	25TB
2.50TB	90%				
...					

2. Identify an aggregate to which you can move the FlexGroup volume constituent: `volume move target-aggr show -vserver svm_name -volume vol_constituent_name`

The available space in the aggregate that you select must be greater than the size of the FlexGroup volume constituent that you are moving.

```
cluster1::> volume move target-aggr show -vserver vs2 -volume fg1_0002
```

Aggregate Name	Available Size	Storage Type
aggr2	467.9TB	hdd
node12a_aggr3	100.34TB	hdd
node12a_aggr2	100.36TB	hdd
node12a_aggr1	100.36TB	hdd
node12a_aggr4	100.36TB	hdd

5 entries were displayed.

3. Verify that the FlexGroup volume constituent can be moved to the intended aggregate: `volume move start -vserver svm_name -volume vol_constituent_name -perform-validation-only true`

```
cluster1::> volume move start -vserver vs2 -volume fg1_0002 -destination
-aggregate node12a_aggr3 -perform-validation-only true
```

Validation succeeded.

4. Move the FlexGroup volume constituent: `volume move start -vserver svm_name -volume vol_constituent_name -destination-aggregate aggr_name [-allow-mixed-aggr-types {true|false}]`

The volume move operation runs as a background process.

Beginning with ONTAP 9.5, you can move FlexGroup volume constituents from a Fabric Pool to a non-Fabric Pool, or vice versa by setting the `-allow-mixed-aggr-types` parameter to `true`. By default, the

-allow-mixed-aggr-types option is set to false.



You cannot use the `volume move` command for enabling encryption on FlexGroup volumes.

```
cluster1::> volume move start -vserver vs2 -volume fg1_002 -destination  
-aggregate node12a_aggr3
```



If the volume move operation fails due to an active SnapMirror operation, you should abort the SnapMirror operation by using the `snapmirror abort -h` command. In some cases, the SnapMirror abort operation might also fail. In such situations, you should abort the volume move operation and retry later.

5. Verify the state of the volume move operation: `volume move show -volume vol_constituent_name`

The following example shows the state of a FlexGroup constituent volume that completed the replication phase and is in the cutover phase of the volume move operation:

```
cluster1::> volume move show -volume fg1_002  
Vserver    Volume      State      Move Phase  Percent-Complete  Time-To-  
Complete  
-----  
vs2        fg1_002     healthy    cutover     -                  -
```

Use aggregates in FabricPool for existing FlexGroup volumes

Beginning with ONTAP 9.5, FabricPool is supported for FlexGroup volumes. If you want to use aggregates in FabricPool for your existing FlexGroup volumes, you can either convert the aggregates on which the FlexGroup volume resides to aggregates in FabricPool or migrate the FlexGroup volume constituents to aggregates in FabricPool.

What you'll need

- The FlexGroup volume must have space-guarantee set to `none`.
- If you want to convert the aggregates on which the FlexGroup volume resides to aggregates in FabricPool, the aggregates must be using all SSD disks.

About this task

If an existing FlexGroup volume resides on non-SSD aggregates, you must migrate the FlexGroup volume constituents to aggregates in FabricPool.

Choices

- To convert the aggregates on which the FlexGroup volume resides to aggregates in FabricPool, perform the following steps:

- a. Set the tiering policy on the existing FlexGroup volume: `volume modify -volume flexgroup_name -tiering-policy [auto|snapshot|none|backup]`

```
cluster-2::> volume modify -volume fg1 -tiering-policy auto
```

- b. Identify the aggregates on which the FlexGroup volume resides: `volume show -volume flexgroup_name -fields aggr-list`

```
cluster-2::> volume show -volume fg1 -fields aggr-list
vserver volume aggr-list
-----
vs1      fg1      aggr1,aggr3
```

- c. Attach an object store to each aggregate listed in the aggregate list: `storage aggregate object-store attach -aggregate aggregate name -name object-store-name -allow -flexgroup true`

You must attach all of the aggregates to an object store.

```
cluster-2::> storage aggregate object-store attach -aggregate aggr1
-object-store-name Amazon01B1
```

- To migrate the FlexGroup volume constituents to aggregates in FabricPool, perform the following steps:

- a. Set the tiering policy on the existing FlexGroup volume: `volume modify -volume flexgroup_name -tiering-policy [auto|snapshot|none|backup]`

```
cluster-2::> volume modify -volume fg1 -tiering-policy auto
```

- b. Move each constituent of the FlexGroup volume to an aggregate in FabricPool in the same cluster: `volume move start -volume constituent-volume -destination-aggregate FabricPool_aggregate -allow-mixed-aggr-types true`

You must move all FlexGroup volume constituents to aggregates in FabricPool (in case the FlexGroup volume constituents are on mixed aggregate types) and ensure that all the constituents are balanced across the nodes in the cluster.

```
cluster-2::> volume move start -volume fg1_001 -destination-aggregate
FP_aggr1 -allow-mixed-aggr-types true
```

Related information

Rebalance FlexGroup volumes

Beginning with ONTAP 9.12.1, you can rebalance FlexGroup volumes by non-disruptively moving files from one constituent in a FlexGroup to another constituent.

FlexGroup rebalancing helps redistribute capacity when imbalances develop over time due to the addition of new files and file growth. After you manually start the rebalance operation, ONTAP selects the files and moves them automatically and non-disruptively.

Automatic rebalancing is available only when all nodes in the cluster are running ONTAP 9.12.1 or later releases. You must enable multipart inode granular data functionality on any FlexGroup volume that runs the rebalancing operation. Once that functionality is enabled, you cannot revert to a previous ONTAP version unless you delete the FlexGroup and restore a previous version.

FlexGroup rebalancing considerations

You should be aware of how FlexGroup rebalancing works and how it interacts with other ONTAP features.

- FlexVol to FlexGroup conversion

It is recommended that you *not* use automatic FlexGroup rebalancing after a FlexVol to FlexGroup conversion. Instead, you can use the disruptive retroactive file move feature available in ONTAP 9.10.1 and later, by entering the `volume rebalance file-move` command. For command syntax, see the `volume rebalance file-move start` man page.

Rebalancing with the non-disruptive retroactive file move feature can degrade performance when moving large numbers of files, like when you perform a FlexVol to FlexGroup conversion, and as much as 50 to 85% of the data on the FlexVol volume is moved to a new constituent.

- Minimum and maximum file size

File selection for automatic rebalancing is based on blocks saved. The minimum file size considered for rebalancing is 100 MB by default (can be configured as low as 4KB using the `min-file-size` parameter shown below) and the maximum file size is 100 GB.

- Files in Snapshot copies

You can configure FlexGroup rebalancing to only consider files to be moved which are not currently present in any Snapshot copies. When rebalancing is started, a notification displays if a Snapshot copy operation is scheduled anytime during a rebalancing operation.

Snapshot copies are restricted if a file is being moved and is undergoing framing at the destination. A Snapshot copy restore operation is not allowed while file rebalancing is in progress.

- SnapMirror operations

FlexGroup rebalancing should take place between scheduled SnapMirror operations. A SnapMirror operation might fail if a file is being relocated before a SnapMirror operation begins if that file move does not complete within the 24-minute SnapMirror retry period. Any new file relocation that begins after a SnapMirror transfer has started will not fail.

- File-based compression storage efficiency

With file-based compression storage efficiency, the file is decompressed before it's moved to the destination, so the compression savings is lost. The compression savings is regained after a manually initiated background scanner runs on the FlexGroup volume after rebalancing. However, if any file is associated with a Snapshot copy on any volume, the file will be ignored for compression.

- Deduplication

Moving deduplicated files can cause increased overall usage for the FlexGroup volume. During file rebalancing, only unique blocks are moved to the destination, freeing that capacity on the source. Shared blocks remain on the source and are copied to the destination. While this achieves the goal of reducing the used capacity on a nearly full source constituent, it can also lead to increased overall usage on the FlexGroup volume due to copies of shared blocks on the new destinations. This is also possible when files that are part of a Snapshot copy are moved. The space savings is not fully recognized until the Snapshot copy schedule recycles and there are no longer copies of the files in Snapshot copies.

- FlexClone volumes

If file rebalancing is in progress when a FlexClone volume is created, the rebalancing will not be performed on the FlexClone volume. Rebalancing on the FlexClone volume should be performed after it is created.

- File move

When a file is moved during a FlexGroup rebalancing operation, the file size is reported as part of quota accounting on both the source and destination constituents. Once the move is completed, quota accounting returns to normal, and the file size is only reported on the new destination.

Enable FlexGroup rebalancing

Beginning with ONTAP 9.12.1, you can enable automatic nondisruptive FlexGroup volume rebalancing to redistribute files between FlexGroup constituents.

Before you begin

You must have enabled the `granular-data` option on the FlexGroup volume before enabling FlexGroup rebalancing. You can enable it by using one of these methods:

- When you create FlexGroup volume using the `volume create` command
- By modifying an existing FlexGroup volume to enable the setting using the `volume modify` command
- Setting it automatically when FlexGroup rebalancing is initiated using the `volume rebalance` command

You can manage FlexGroup rebalancing by using ONTAP System Manager or the ONTAP CLI.

System Manager

Steps

1. Navigate to **Storage > Volumes** and locate the FlexGroup volume to rebalance.
2. Click  to view the volume details.
3. Click **Rebalance**.

CLI

Steps

1. Start automatic rebalancing: `volume rebalance start -vserver SVM_name -volume volume_name`

Optionally, you can specify the following options: `[-max-runtime] <time interval>` Maximum Runtime `[-max-threshold <percent>]` Maximum Imbalance Threshold per Constituent `[-min-threshold <percent>]` Minimum Imbalance Threshold per Constituent `[-max-file-moves <integer>]` Maximum Concurrent File Moves per Constituent `[-min-file-size {<integer>[KB|MB|GB|TB|PB]}]` Minimum file size `[-exclude-snapshots {true|false}]` Exclude files stuck in Snapshot copies

Example:

```
volume rebalance start -vserver vs0 -volume fg1
```

2. In the **Rebalance Volume** window, change the default settings if necessary.
3. Click **Save** to accept the default settings.

Modify FlexGroup rebalance configurations

You can change a FlexGroup rebalancing configuration to update the imbalance threshold, number of concurrent files moves minimum file size, maximum runtime, and to include or exclude Snapshot copies.

System Manager

Steps

1. Navigate to **Storage > Volumes** and locate the FlexGroup volume to rebalance.
2. Click  to view the volume details.
3. Click **Rebalance**.
4. In the **Rebalance Volume** window, change the default settings as needed.
5. Click **Save**.

CLI

Step

1. Modify automatic rebalancing: `volume rebalance modify -vserver SVM_name -volume volume_name`

You can specify one or more of the following options: `[-max-runtime] <time interval>` Maximum Runtime `[-max-threshold <percent>]` Maximum Imbalance Threshold per Constituent `[-min-threshold <percent>]` Minimum Imbalance Threshold per Constituent `[-max-file-moves <integer>]` Maximum Concurrent File Moves per Constituent `[-min-file-size {<integer>[KB|MB|GB|TB|PB]}]` Minimum file size `[-exclude-snapshots {true|false}]` Exclude files stuck in Snapshot copies

Stop FlexGroup rebalance

After FlexGroup rebalancing is enabled, you can stop it at any time.

System Manager

Steps

1. Navigate to **Storage > Volumes** and locate the FlexGroup volume.
2. Click  to view the volume details.
3. Click **Stop Rebalance**.

CLI

Step

1. Stop FlexGroup rebalancing: `volume rebalance stop -vserver SVM_name -volume volume_name`

View FlexGroup rebalance status

You can display the status about a FlexGroup rebalance operation, the FlexGroup rebalance configuration, the rebalance operation time, and the rebalance instance details.

System Manager

Steps

1. Navigate to **Storage > Volumes** and locate the FlexGroup volume.
2. Click  to view the FlexGroup details.
3. **FlexGroup Balance Status** is displayed near the bottom of the details pane.
4. To view information about the last rebalance operation, click **Last Volume Rebalance Status**.

CLI

Step

1. View the status of a FlexGroup rebalance operation: `volume rebalance show`

Example of rebalance state:

```
> volume rebalance show
Vserver: vs0
```

Imbalance					Target
Volume	State	Total	Used	Used	
Size	%				
fg1	idle	4GB	115.3MB	-	
8KB	0%				

Example of rebalance configuration details:

```
> volume rebalance show -config
Vserver: vs0
```

Max		Threshold		Max
Min	Exclude	Min	Max	File Moves
Volume	Runtime			
File Size	Snapshot			
fg1	6h0m0s	5%	20%	25
4KB	true			

Example of rebalance time details:

```
> volume rebalance show -time
Vserver: vs0
Volume                Start Time                Runtime
Max Runtime
-----
fgl                    Wed Jul 20 16:06:11 2022    0h1m16s
6h0m0s
```

Example of rebalance instance details:

```
> volume rebalance show -instance
Vserver Name: vs0
Volume Name: fgl
Is Constituent: false
Rebalance State: idle
Rebalance Notice Messages: -
Total Size: 4GB
AFS Used Size: 115.3MB
Constituent Target Used Size: -
Imbalance Size: 8KB
Imbalance Percentage: 0%
Moved Data Size: -
Maximum Constituent Imbalance Percentage: 1%
Rebalance Start Time: Wed Jul 20 16:06:11 2022
Rebalance Stop Time: -
Rebalance Runtime: 0h1m32s
Rebalance Maximum Runtime: 6h0m0s
Maximum Imbalance Threshold per Constituent: 20%
Minimum Imbalance Threshold per Constituent: 5%
Maximum Concurrent File Moves per Constituent: 25
Minimum File Size: 4KB
Exclude Files Stuck in Snapshot Copies: true
```

Copyright information

Copyright © 2023 NetApp, Inc. All Rights Reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means—graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system—without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP “AS IS” AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

LIMITED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (b)(3) of the Rights in Technical Data -Noncommercial Items at DFARS 252.227-7013 (FEB 2014) and FAR 52.227-19 (DEC 2007).

Data contained herein pertains to a commercial product and/or commercial service (as defined in FAR 2.101) and is proprietary to NetApp, Inc. All NetApp technical data and computer software provided under this Agreement is commercial in nature and developed solely at private expense. The U.S. Government has a non-exclusive, non-transferrable, nonsublicensable, worldwide, limited irrevocable license to use the Data only in connection with and in support of the U.S. Government contract under which the Data was delivered. Except as provided herein, the Data may not be used, disclosed, reproduced, modified, performed, or displayed without the prior written approval of NetApp, Inc. United States Government license rights for the Department of Defense are limited to those rights identified in DFARS clause 252.227-7015(b) (FEB 2014).

Trademark information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.